NILADRISH CHATTERJEE

Computer Architect

@ nilcsutah@gmail.com 801-554-1359 Kirkland, WA www.niladrish.org Google Scholar

OVERVIEW

Computer architect with 10+years of industry and academic research experience. At NVIDIA research, I develop proof-of-concepts of next-generation memory and system architectures to accelerate HPC and AI workloads. Some of my notable recent contributions at NVIDIA include

- post-HBM high-bandwidth on-package memory architecture
- highly energy-efficient (less than 2pJ/b) DRAM for irregular workloads
- reimagined GPU memory hierarchy for additional capacity at high bandwidth
- customized GPU architecture for AI vs HPC architectures
- fast and silicon-accurate GPU simulator with a detailed memory-controller model for architecture exploration

My research has influenced the technical roadmap at NVIDIA and of partner memory vendors for future products. My research has been published in top computer architecture conferences, covered by the tech media and led to patents. I regularly serve on the program committee of conferences and as reviewer of journals on computer architecture.

PUBLICATIONS

Memory Systems Architecture

- **HPCA-2022** Saving PAM4 Bus Energy with SMOREs: Sparse Multi-level Opportunistic Restricted Encodings.
  M. O'Connor, D. Lee, N. Chatterjee, M. Sullivan and S. W. Keckler.
- **ACM TACO** GPU Domain Specialization via Composable On-Package Architecture.
  Y. Fu, E. Bolotin, N. Chatterjee, D. Nellans and S. W. Keckler.
- **HPCA-2018** Reducing Data Transfer Energy by Exploiting Similarity within a Data Transaction.
  D. Lee, M. O'Connor, and N. Chatterjee.
- **MICRO-2017** Fine-Grained DRAM: Energy-Efficient DRAM for Extreme Bandwidth Systems.
  M. O'Connor, N. Chatterjee, D. Lee, J. Wilson, A. Agrawal, S. W. Keckler, and W. J. Dally.
- **HPCA-2017** Architecting an Energy-Efficient Memory System for GPUs.
  N. Chatterjee, M. O'Connor, D. Lee, D. R. Johnson, M. Rhu, S. W. Keckler, and W. J. Dally.
- **ISPASS-2016** Addressing Service Interruptions in Memory with Thread-to-Rank Assignment.
  M. Shevgoor, R. Balasubramonian, N. Chatterjee, and J. Kim.
- **SC-2014** Managing DRAM Latency Divergence in Irregular GPGPU Applications.
  N. Chatterjee, M. O'Connor, G. H. Loh, N. Jayasena, and R. Balasubramonian.
- **MICRO-2013** Quantifying the Relationship between the Power Delivery Network and Architectural Policies in 3D-Stacked Memory Devices.
  M. Shevgoor, J.-S. Kim, N. Chatterjee, R. Balasubramonian, A. Davis, and A. N. Udipi.
- **MICRO-2012** Leveraging Heterogeneity in DRAM Main Memories to Accelerate Critical Word Access.

EXPERIENCE

NVIDIA

Sr. Research Scientist

2013 – Present Redmond, WA

- Researching fundamental advances in memory systems for HPC and AI platforms.
- Co-inventor of energy-efficient, high-performance DRAM substrate (FGDRAM) for post-HBM systems
- Performance modeling and simulation
- Influencing GPU and memory roadmaps and architectural features

HP Labs

Intern

Fall 2012 SLC, UT

- Disaggregated memory systems for high-capacity datacenters.

AMD Research

Intern

Spring 2012 Sunnyvale, CA

- Memory scheduling algorithms for multi-client SoC memory controllers.
- SIMT-aware memory controller to minimize latency-divergence while balancing sustained throughput.

EDUCATION

Ph.D. in Computer Engineering

University of Utah

2008 – 2013

Designing Efficient Memory Schedulers For Future Systems

Advisor: Dr. Rajeev Balasubramonian

B.E. in Computer Science

Jadavpur University

2003 – 2007

N. Chatterjee, M. Shevgoor, R. Balasubramonian, A. Davis, Z. Fang, R. Illikkal, and R. Iyer.

- **HPCA-2012** Staged-Reads: Mitigating the Impact of DRAM Writes on DRAM Reads.
  N. Chatterjee, R. Balasubramonian, N. Muralimanohar, A. Davis, and N. Jouppi.
- **ISCA-2010** Rethinking DRAM Design and Organization for Energy-Constrained Multi-cores.
  A. Udipi, N. Muralimanohar, N. Chatterjee, R. Balasubramonian, A. Davis, and N. Jouppi.
- **ASPLOS-2010** Micro-pages: Increasing DRAM Efficiency with Locality-Aware Data Placement.
  K. Sudan, N. Chatterjee, M. Awasthi, D. Nellans, R. Balasubramonian, and A. Davis.

Hardware Architectures for Deep Learning

- **ISPASS-2021** Learning Sparse Matrix Row Permutations for Efficient SpMM on GPU Architectures.
  A. Mehrabi, D. Lee, N. Chatterjee, D. Sorin, B. Lee, M. O'Connor.
- **ISPASS-2019** DeLTA: GPU Performance Model for Deep Learning Applications with In-Depth Memory System Traffic Analysis.
  S. Lym, D. Lee, M. O'Connor, N. Chatterjee, M. Erez.
- **HPCA-2018** Compressing DMA Engine: Leveraging Activation Sparsity for Training Deep Neural Networks.
  M. Rhu, M. O'Connor, N. Chatterjee, J. Pool, Y. Kwon, and S. W. Keckler.

Near-Memory Data Processing

- **SC-2019** Near-Memory Data Transformation for Efficient Sparse Matrix Multi-Vector Multiplication.
  D. Fujiki, N. Chatterjee, M. O'Connor, and D. Lee.
- **SC-2017** Toward Standardized Near-Data Processing with Unrestricted Data Placement for GPUs.
  G. Kim, N. Chatterjee, M. O'Connor, and K. Hsieh.
- **ISCA-2016** Transparent Offloading and Mapping (TOM): Enabling Programmer-Transparent Near-Data Processing in GPU Systems.
  K. Hsieh, E. Ebrahimi, G. Kim, N. Chatterjee, M. O'Connor, N. Vijaykumar, O. Mutlu, and S. W. Keckler.

Modeling and Simulation

- **HPCA-2021** Need for Speed: Experiences Building a Trustworthy System-Level GPU Simulator.
  O. Villa, D. Lustig, Z. Yan, E. Bolotin, N. Chatterjee, N. Jiang, D. Nellans.
- **UUCS-TR-12-002** USIMM: the Utah Simulated Memory Module. A Simulation Infrastructure for the JWAC Memory Scheduling Competition.
  N. Chatterjee, R. Balasubramonian, M. Shevgoor, S. H. Pugsley, A. N. Udipi, A. Shaifei, M. Awasthi, K. Sudan, and Z. Chishti
- **SIGMETRICS-2018** What Your DRAM Power Models are Not Telling You: Lessons from a Detailed Experimental Study .
  S. Ghose, A. G. Yaglicki, R. Gupta, D. Lee, K. Kudrolli, W. X. Liu, H. Hassan, K. K. Chang, N. Chatterjee, A. Agrawal, M. O'Connor, and O. Mutlu .
- **SIGMETRICS-2017** Understanding Reduced-Voltage Operation in Modern DRAM Devices: Experimental Characterization, Analysis, and Mechanisms.
  K. K. Chang, A. G. Yaglikci, S. Ghose, A. Agrawal, N. Chatterjee, A. Kashyap, D. Lee, M. O'Connor, H. Hassan, and O. Mutlu.

ISSUED PATENTS

10,468,093 Systems and methods for dynamic random access memory (DRAM) sub-channels

9,910,605 Page migration in a hybrid memory device

9,846,550 Memory access methods and apparatus (divisional of 9,361,955)

9,535,831 Page migration in a 3D stacked hybrid memory

9,489,321 Scheduling memory accesses using an efficient row burst value

9,361,955 Memory Access Methods and Apparatus

HONORS

Best Paper Award
ISPASS 2016

Best Paper Nominee
HPCA 2018

Best Poster Awards
HiPC 2009. School of Computing Graduate Research Competition 2010 and 2013

School of Computing Fellowship
University of Utah teaching fellowship 2008. Awarded to outstanding incoming graduate students.

Media Coverage
HPCA 2017 work on GPU DRAM microarchitecture covered on The Next Platform

Media Coverage
ACM TACO work on Composable GPU Architectures covered on The Next Platform

SERVICE

Conference Program Committee Member
HPCA 2019, HPCA 2018, IPDPS 2017, ICS 2016

External Review Committee Member
MICRO 2022/2021/2019/2012, HPCA 2021/2015/2012, ISCA 2021/2015/2014, HPG 2017

Journal Reviewer
ACM TACO

Co-organizer
3rd JILP Workshop on Computer Architecture Competitions (Memory Scheduling Championship). Held with ISCA-2012.